

# Chapter 4: Parallel Program Structures I

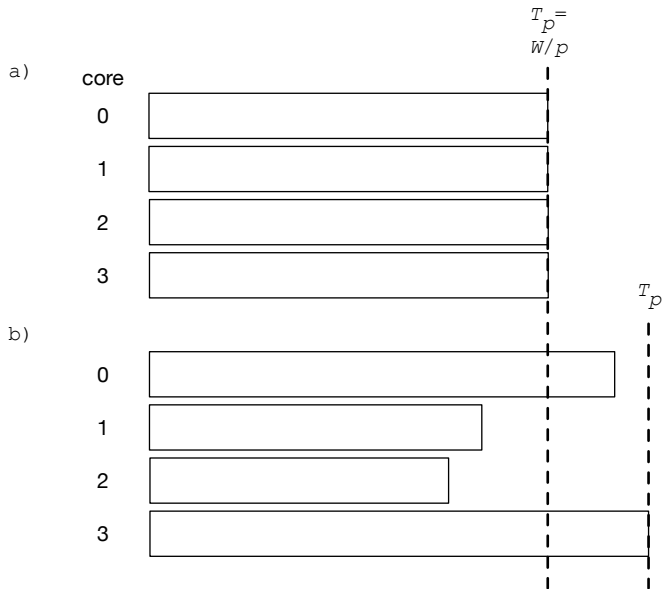
Elements of Parallel Computing

Eric Aubanel

# Wall clock time

- ▶ start timer as soon as program launched
- ▶ stop timer once all cores finish execution

# Load Balance



# SIMD: Strictly Data Parallel

SIMD execution can be enabled via:

- ▶ automatic compiler vectorization
  - ▶ better if loops annotated by programmer (pragmas)
- ▶ array notation:
  - ▶  $c[0 : n - 1] \leftarrow a[0 : n - 1] + b[0 : n - 1]$
  - ▶  $c \leftarrow a + b$
- ▶ forall loops

# Strip-Mining

Compiler break down iterations into chunks with a number of elements that depends on the number of bits in the data type

E.g. assume 512 bit vector registers, double precision arithmetic:  $2^{20}$  operations broken into  $2^{17}$  groups of 8

# SIMD Notation

We will use set notation:

$$\{c[i] \leftarrow a[i] + b[i] : i \in [0..n)\}$$

# Row-wise SIMD Matrix-Vector Multiplication

```
 $\{b[i] \leftarrow 0 : i \in [0..n)\}$   
for  $j \leftarrow 0$  to  $m - 1$  do  
     $\{b[i] \leftarrow b[i] + A[i, j] * x[j] : i \in [0..n)\}$   
end
```


$$\begin{bmatrix} A_{00} & A_{01} & A_{02} & A_{03} \\ A_{10} & A_{11} & A_{12} & A_{13} \\ A_{20} & A_{21} & A_{22} & A_{23} \\ A_{30} & A_{31} & A_{32} & A_{33} \end{bmatrix} \times \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} A_{00}x_0 \\ A_{10}x_0 \\ A_{20}x_0 \\ A_{30}x_0 \end{bmatrix}$$

$$\begin{bmatrix} A_{00} & A_{01} & A_{02} & A_{03} \\ A_{10} & A_{11} & A_{12} & A_{13} \\ A_{20} & A_{21} & A_{22} & A_{23} \\ A_{30} & A_{31} & A_{32} & A_{33} \end{bmatrix} \times \begin{bmatrix} x_0 \\ x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} A_{01}x_1 \\ A_{11}x_1 \\ A_{21}x_1 \\ A_{31}x_1 \end{bmatrix}$$



# Row-wise SIMD Matrix-Vector Multiplication

- Requires loading of column of matrix
  - not efficient if stored in row-major order



# Column-wise Matrix-Vector Multiplication

Inner product of each row of  $A$  with  $b$  in parallel:  
**reduction**

# Reduction

Use SIMD *conditional execution*:

```
for  $k \leftarrow 0$  to  $\log n - 1$  do  
     $j \leftarrow 2^k$   
     $\{a[i] \leftarrow a[i] + a[i + j] : i \in [0..n) \mid i \bmod 2j = 0\}$   
end  
// result in  $a[0]$ 
```

# Control Divergence

Conditional execution can reduce performance

E.g. say SIMD width stores 4 elements of  $a$

- ▶ first stage of reduction ( $k = 0$ ): only two additions could be done in parallel
- ▶ other stages: additions would all be serialized.

# Change Indexing





```
for  $k \leftarrow 0$  to  $\log n - 1$  do  
     $j \leftarrow 2^{k+1}$   
     $\{a[i * j] \leftarrow a[i * j] + a[i * j + j/2] : i \in [0..n/j)\}$   
end
```

...but has scattered memory access pattern

# Divergence-Free Reduction

Change order of additions:

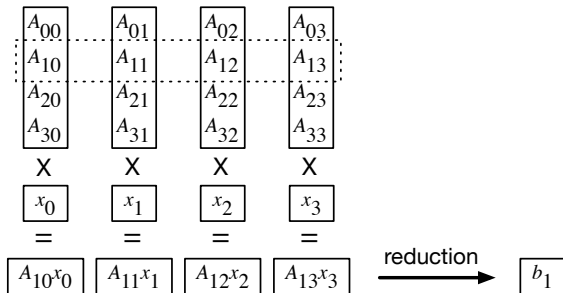
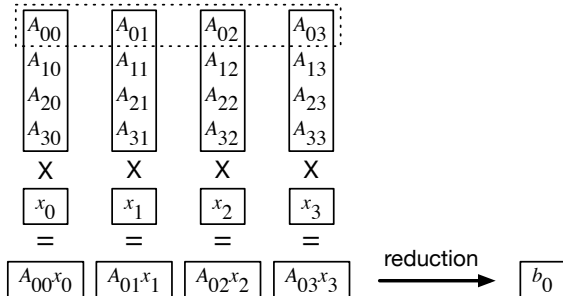
```
for  $k \leftarrow \log n - 1$  to 0 do  
     $j \leftarrow 2^k$   
     $\{a[i] \leftarrow a[i] + a[i + j] : i \in [0..j)\}$   
end  
// result in  $a[0]$ 
```



# Column-wise Matrix-Vector Multiplication

```
{  $b[i] \leftarrow 0 : i \in [0..n)$  }  
for  $i \leftarrow 0$  to  $n - 1$  do  
    {  $temp[j] \leftarrow A[i, j] * x[j] : j \in [0..m)$  }  
    for  $k \leftarrow \log n - 1$  to  $0$  do  
        {  $temp[j] \leftarrow temp[j] + temp[j + 2^k] : j \in$   
           $[0..2^k)$  }  
    end  
     $b[i] \leftarrow temp[0]$   
end
```





# SIMD Subset Sum

**for**  $i \leftarrow 2$  *to*  $n$  **do**

$\{F[i, j] \leftarrow F[i - 1, j] : j \in [1..S]\}$

$\{F[i, j] \leftarrow F[i, j] \vee F[i - 1, j - s[i]] : j \in [1..S] \mid$   
 $j \geq s[i]\}$

**end**

Control divergence only for chunks where  $j < s[i]$   
and  $j \geq s[i]$

# SIMD Guidelines

1. Watch out for dependencies
2. Avoid control divergence
3. Optimize memory access patterns

# Shared Memory Programming

- ▶ Fork-Join
- ▶ Parallel loops
- ▶ Tasks with dependencies
- ▶ Single Program Multiple Data (SPMD)

# Threads

Thread:

- ▶ has its own program counter and private memory region in its stack frame
- ▶ shares instructions, heap and global memory with other threads.
- ▶ has a thread id

# Fork-Join

- ▶ Child tasks are *forked* by parent
- ▶ Parents may need to wait for children to complete
  - ▶ called *joining*
- ▶ Mainly used by recursively creating tasks

# Recursive Estimation of Pi

```
//  $n = 2^k$  experiments
```

```
Procedure estimatePi( $n$ )
```

```
    return recPi( $n$ )*4/ $n$ 
```

```
end
```

```
// returns sum of points in circle
```

```
Procedure recPi( $n$ )
```

```
    if  $n = 1$  then
```

```
         $sum \leftarrow 0$ 
```

```
         $x \leftarrow$  pseudo-random number  $\in [-1, 1]$ 
```

```
         $y \leftarrow$  pseudo-random number  $\in [-1, 1]$ 
```

```
        if  $x^2 + y^2 \leq 1$  then  $sum \leftarrow 1$ 
```

```
    else
```


```
         $sum \leftarrow$  recPi( $n/2$ ) + recPi( $n/2$ )
```

```
    end
```

```
    return  $sum$ 
```

```
end
```

# Reduction Task Graph Mapped to Threads





# Fork Tasks not Threads

Fork and join tasks and let runtime system assign them to threads.

- ▶ **spawn** = fork
- ▶ **sync** = join

$sum1 \leftarrow \mathbf{spawn} \text{ recPi}(n/2)$

$sum2 \leftarrow \text{recPi}(n/2)$

**sync**

$sum \leftarrow sum1 + sum2$

# Sequential Cutoff

- ▶ Too few tasks: difficult to balance work across threads
- ▶ Too many tasks: significant runtime overhead

Choose base case to limit number of tasks

```

Procedure recPi( $n$ )
  if  $n < \text{cutoff}$  then
     $sum \leftarrow 0$ 
    for  $i \leftarrow 1$  to  $n$  do
       $x \leftarrow$  pseudo-random number  $\in [-1, 1]$ 
       $y \leftarrow$  pseudo-random number  $\in [-1, 1]$ 
      if  $x^2 + y^2 \leq 1$  then  $sum \leftarrow sum + 1$ 
    end
  else
     $sum1 \leftarrow$  spawn recPi( $n/2$ )
     $sum2 \leftarrow$  recPi( $n/2$ )
    sync
     $sum \leftarrow sum1 + sum2$ 
  end
  return  $sum$ 
end

```

# Merge Sort

- ▶ Use fork-join for both splitting and merging
- ▶ Also use sequential cutoffs for both
- ▶ split the merge:
  - ▶ find the median *mid* of longest of the sorted subarrays
  - ▶ binary search of the other subarray to find the index that splits it into elements less than and greater than *mid*

**Input:** array  $a$  of length  $n$ .

**Output:** array  $b$ , containing array  $a$  sorted. Array  $a$  overwritten.

```
arrayCopy( $a$ ,  $b$ ) // copy array  $a$  to  $b$ 
```

```
parMergeSort( $a$ , 0,  $n$ ,  $b$ )
```

```
// sort elements with index  $i \in [lower..upper)$ 
```

```
Procedure parMergeSort( $a$ ,  $lower$ ,  $upper$ ,  $b$ )
```

```
    if ( $upper - lower$ ) < cutoff then
```

```
        sequentialSort( $a$ ,  $lower$ ,  $upper$ ,  $b$ )
```

```
    else
```

```
         $mid \leftarrow \lfloor (upper + lower)/2 \rfloor$ 
```

```
        spawn parMergeSort( $b$ ,  $lower$ ,  $mid$ ,  $a$ )
```

```
        parMergeSort( $b$ ,  $mid$ ,  $upper$ ,  $a$ )
```

```
    sync
```

```
    parMerge( $a$ ,  $lower$ ,  $mid$ ,  $mid$ ,  $upper$ ,  $b$ ,  $lower$ )
```

```
    return
```

```
end
```

```
end
```

# Parallel Merge

- ▶ parallel merge of sorted sub-arrays  
 $i \in [low1..up1)$  and  $i \in [low2..up2)$  of  $a$  into  $b$   
starting at index  $start$
- ▶ two indices for each of lower and upper  
subarrays, since they won't always be contiguous

```

Procedure parMerge(a, low1, up1, low2, up2, b, start)
     $k1 \leftarrow up1 - low1, k2 \leftarrow up2 - low2$ 
    if  $k1 + k2 < \text{cutoff}$  then
        sequentialMerge(a, low1, up1, low2, up2, b, start)
    else
        ...
    end
end

```

```

if  $k1 + k2 < \text{cutoff}$  then
    sequentialMerge(a, low1, up1, low2, up2, b, start)
else
    if  $k1 \geq k2$  then
         $\text{mid1} \leftarrow \lfloor (\text{low1} + \text{up1} - 1) / 2 \rfloor$ 
        // mid2: first index in [low2, up2) such
           that  $a[\text{index}] > a[\text{mid1}]$ 
         $\text{mid2} \leftarrow \text{binarySearch}(\text{a}, \text{low2}, \text{up2}, \text{mid1})$ 
    else
         $\text{mid2} \leftarrow \lfloor (\text{low2} + \text{up2} - 1) / 2 \rfloor$ 
         $\text{mid1} \leftarrow \text{binarySearch}(\text{a}, \text{low1}, \text{up1}, \text{mid2}) - 1$ 
         $\text{mid2} \leftarrow \text{mid2} + 1$ 
    end
    spawn parMerge(a, low1, mid1 + 1, low2, mid2, b,
        start)
    parMerge(a, mid1 + 1, up1, mid2, up2, b,
        start + mid1 - low1 + 1 + mid2 - low2)
    sync
end

```



# Fork-Join Guidelines

- ▶ Don't create more independent tasks than necessary
- ▶ Use a sequential cutoff to limit the depth of the recursion
- ▶ Avoid unnecessary allocation of memory
- ▶ Be careful if shared read/write access to data required
  - ▶ **Warning:** may be hidden inside functions