

Chapter 4: Parallel Program Structures IV

Elements of Parallel Computing

Eric Aubanel

Master-Worker

- ▶ Good for load balancing
- ▶ Master queues tasks (may be implicit, i.e., tasks known)
- ▶ Workers de-queue tasks
- ▶ Master can also act as worker
- ▶ *distributed memory*: master sends tasks and gathers results
- ▶ *shared memory*: need to avoid race conditions, for conflicts between master and workers or between workers

Shared Memory Master-Worker Fractal

```
shared kount, chunkCount
if id = 0 then
    chunkCount  $\leftarrow$  nt * chunk
end
istart  $\leftarrow$  id * chunk
iend  $\leftarrow$  (id + 1) * chunk - 1
barrier()
while istart  $\leq$  n - 1 do
    for i  $\leftarrow$  istart to iend do
        ...// see sequential Algorithm ??
    end
    begin critical
        istart  $\leftarrow$  chunkCount
        chunkCount  $\leftarrow$  chunkCount + chunk
    end critical
    iend  $\leftarrow$  min(istart + chunk - 1, n - 1)
end
```

Master-Worker Considerations

- ▶ Need to choose good task (chunk) size for workers
- ▶ Master can get overloaded
- ▶ Alternatives include multiple masters, work stealing

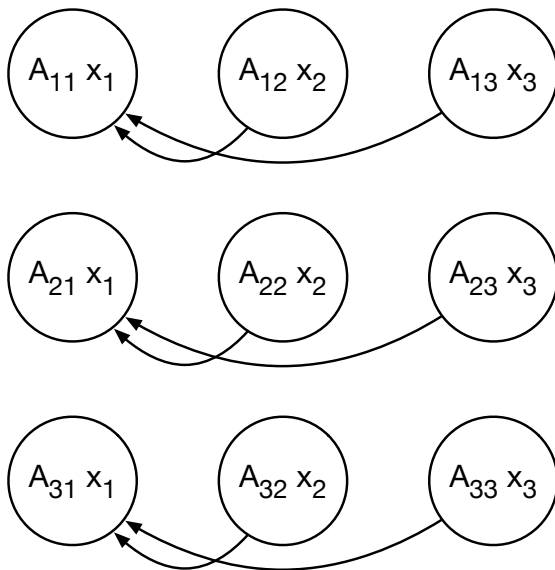
Distributed Memory Programming

- ▶ Distributed arrays
- ▶ Message passing
- ▶ Local and global communication


Distributed Arrays: Matrix-Vector Multiplication

$$\begin{bmatrix} A_{11} & A_{12} & A_{13} \\ A_{21} & A_{22} & A_{23} \\ A_{31} & A_{32} & A_{33} \end{bmatrix} \times \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} A_{11}x_1 \\ A_{21}x_1 \\ A_{31}x_1 \end{bmatrix} + \begin{bmatrix} A_{12}x_2 \\ A_{22}x_2 \\ A_{32}x_2 \end{bmatrix} + \begin{bmatrix} A_{13}x_3 \\ A_{23}x_3 \\ A_{33}x_3 \end{bmatrix}$$

Task Graph



Redistribution of Results



Message Passing

SPMD, with local and global communication

Local Communication

Two-sided: both sender and receiver participate.
E.g., exchanging values:

```
if id = 0 then  
    nonblocking send data to 1  
    receive data from 1  
else  
    nonblocking send data to 0  
    receive data from 0  
end
```

Will this work?

Local Communication

Send/receive: blocking/non-blocking,
synchronous/asynchronous

- ▶ blocking:
 - ▶ synchronous: waits until message received
 - ▶ asynchronous: waits until send buffer can be overwritten
- ▶ non-blocking: don't wait. Need method to determine when

Potential Deadlock

if $id = 0$ **then**

blocking send data to 1
 receive data from 1


else

blocking send data to 0
 receive data from 0

end

Deadlock if communication synchronous.

Game of Life



// each task has $(n/p + 2) \times n$ arrays *grid* and *newGrid*

Input: $n \times n$ grid of cells, each with a state of alive (1) or dead (0).

Output: evolution of grid for a given number of generations

$nbDown \leftarrow (id + 1) \bmod p$, $nbUp \leftarrow (id - 1 + p) \bmod p$

$m \leftarrow n/p$ // Assume $n \bmod p = 0$

for *a number of generations* **do**

 // nonblocking send of boundary values to
 neighbors

 nonblocking send *grid*[$m, 0..n - 1$] to *nbDown*

 nonblocking send *grid*[$1, 0..n - 1$] to *nbUp*

 // receive boundary values from neighbors into
 ghost elements


 receive from *nbDown* into *grid*[$m + 1, 0..n - 1$]

 receive from *nbUp* into *grid*[$0, 0..n - 1$]

foreach *cell at coordinate* $(i, j) \in (1..m, 0..n)$ **do**
 updateGridCell(*grid*, *newGrid*, *i*, *j*)

end

Global Communication



Scatter/Gather

- ▶ `scatter(0,a,2,b)`: source task, source data, number of elements, destination data
- ▶ `gather(0,b,2,a)`: destination task, source data, number of elements, destination data

Improving Game of Life

Reading and scattering grid for Game of Life:

```
if  $id = 0$  then  
    read grid from disk  
end  
scatter(0, grid,  $n * n / p$ , grid[1.. $n / p$ , 0.. $n - 1$ ])
```

Gathering and displaying grid:

```
if number of generations mod  $d = 0$  then  
    gather(0, grid[1.. $n / p$ , 0.. $n - 1$ ],  $n * n / p$ , dgrid)  
    if  $id = 0$  then  
        display dgrid  
    end  
end
```

Broadcast/Reduction

- ▶ `broadcast(0,c,size,d)`: source task, source data, size of data, destination data
- ▶ `reduce(0,s,size,op,t)`: destination task, source data, size of data, destination data

Can perform global communication over subset of tasks

- ▶ `broadcast(0,c,size,d, group)`
- ▶ `reduce(0,s,size,op,t, group)`

Row-Wise Matrix-Vector Multiplication

```
 $nb \leftarrow n/p$  // assume  $n \bmod p = 0$   
for  $i \leftarrow 0$  to  $nb - 1$  do  
     $c[i] \leftarrow 0$   
    foreach column  $j$  of  $a$  do  
         $c[i] \leftarrow c[i] + a[i, j] * b[j]$   
    end  
end  
 $\text{gather}(0, c, nb, c)$   
 $\text{broadcast}(0, c, n, c)$ 
```

2D Matrix-Vector Multiplication

// Assume $p = q^2$, matrix is square, $n \bmod q = 0$

$q \leftarrow \sqrt{p}$, $nb \leftarrow n/q$

for $i \leftarrow 0$ to $n/q - 1$ **do**

$c[i] \leftarrow 0$

for $j \leftarrow 0$ to $n/q - 1$ **do**

$c[i] \leftarrow c[i] + a[i, j] * b[j]$

end

end

$rowID \leftarrow \lfloor id/q \rfloor$, $colID \leftarrow id \bmod q$, $destID \leftarrow rowID * q$

$group \leftarrow [destID..destID + q - 1]$

$reduce(destID, c, nb, sum, c, group)$

$sourceID \leftarrow colID * q$


$group \leftarrow [sourceID, colID, colID + q, .., colID + (q - 1) * q]$

$broadcast(sourceID, c, nb, c, group)$

$group \leftarrow [0, q, .., (q - 1) * q]$

$broadcast(0, c, nb, c, group)$

Subset Sum



Input: Array $s[1..n]$ of n positive integers, target sum S

Output: Completed array F

Data: Array $F[1..n, 0..nb - 1]$ initialized to 0 (nb is number of columns owned by task), array $L[0..\lceil S/p \rceil - 1]$ to store received messages

$\text{broadcast}(0, s, n, s)$

$\text{myFirst} \leftarrow \lfloor id * S/p \rfloor + 1$

$\text{myLast} \leftarrow \lfloor (id + 1) * S/p \rfloor$

if $id = 0$ **then** //first block has one extra value

$\text{myFirst} \leftarrow 0$

$F[1, 0] \leftarrow 1$

end

$nb \leftarrow \text{myLast} - \text{myFirst} + 1$

$nL \leftarrow \lceil S/p \rceil$

if $id = \text{findID}(s[1] - 1, p, S)$ **then**

$F[1, s[1] - \text{myFirst}] \leftarrow 1$

```

for  $i \leftarrow 2$  to  $n$  do
     $id1 \leftarrow \text{findID}(\text{myFirst} + s[i] - 1, p, S)$ 
     $id2 \leftarrow \text{findID}(\text{myLast} + s[i] - 1, p, S)$ 
    if  $id1 < p$  then
        // send ...
    end
    if  $id > 0 \wedge \text{myLast} - s[i] \geq 0$  then
        // receive ...
    end
     $\text{solveRow}(F, L, s, nb, \text{myFirst}, i)$ 
end

// Trace  $F$  to return subset, or have task  $p - 1$ 
print  $F[n, \text{myLast}]$  to return yes/no

```

Send

```
if  $id1 < p$  then  
  if  $id1 = id2$  then  
     $myLocalBegin \leftarrow$   
     $\max(0, \lfloor id1 * S/p \rfloor + 1 - s[i] - myFirst)$   
    send  $F[i - 1, myLocalBegin..nb - 1]$  to  $id1$   
  else  
     $destBegin \leftarrow myFirst + s[i]$   
     $destLast \leftarrow \lfloor (id1 + 1) * S/p \rfloor$   
     $nb1 \leftarrow destLast - destBegin + 1$  // # of elements  
    to send to  $id1$   
    if  $id1 > id$  then send  $F[i - 1, 0..nb1 - 1]$  to  $id1$   
    if  $id2 < p$  then send  $F[i - 1, nb1..nb - 1]$  to  $id2$   
  end  
end
```


Receive

```
if  $id > 0 \wedge myLast - s[i] \geq 0$  then  
     $id1 \leftarrow findID(myFirst - s[i] - 1, p, S)$   
    if  $myFirst - s[i] < 0$  then  
         $myLocalBegin \leftarrow s[i] - myFirst$   
    else  
         $myLocalBegin \leftarrow 0$   
    end  
    receive from  $id1$  into  $L[myLocalBegin..nL - 1]$   
     $nS \leftarrow$  size of message received  
     $id2 \leftarrow findID(myLast - s[i] - 1, p, S)$   
    if  $id1 \neq id2 \wedge id2 < id$  then receive from  $id2$  into  
         $L[nS..nL - 1]$   
end
```

```
// return rank of task that owns column  $j$  of array  
// of length  $n$  in a decomposition into  $p$  blocks  
Procedure findID( $j, p, n$ )  
    return  $\lfloor (p * (j + 1) - 1) / n \rfloor$   
end
```

// Solve row i of F with array L for values
needed from other tasks

Procedure solveRow($F, L, s, nb, myFirst, i$)

if $id = 0$ **then**

$F[i, 0] \leftarrow 1$

$jstart \leftarrow 1$

else

$jstart \leftarrow 0$

end

for $j \leftarrow jstart$ **to** $nb - 1$ **do**

$F[i, j] \leftarrow F[i - 1, j]$

$offset \leftarrow j - s[i]$

if $offset \geq 0$ **then** //dependency in my array F

$F[i, j] \leftarrow F[i, j] \vee F[i - 1, offset]$

else if $offset + myFirst \geq 0$ **then** //dependency in
array L

$F[i, j] \leftarrow F[i, j] \vee L[j]$

end

end