

Chapter 5: Performance Analysis

Elements of Parallel Computing

Eric Aubanel

Work-Depth Analysis

- ▶ **Depth:** parallel execution time
- ▶ **Work:** sequential execution time to complete all tasks
- ▶ **Parallelism:** work/depth. Upper bound on performance gain.

Depth

- ▶ Consider all possible paths in task graph
- ▶ Depth given as longest time of all paths
- ▶ If all tasks take the same time then depth is time of longest path: *critical path*
- ▶ No matter how many computational resources, solution cannot be computed in less time than depth

Work Efficiency

Parallel algorithm is **work-efficient** if its work has the same complexity as the best sequential algorithm

Reduction

- ▶ work: $O(n)$
 - ▶ work efficient
- ▶ depth: $O(\log n)$
- ▶ parallelism: $O(n / \log n)$

Naive Merge Sort

- ▶ work: $O(n \log n)$
- ▶ depth: from recurrence relation
 $d(n) = d(n/2) + O(n)$, gives depth of $O(n)$
- ▶ parallelism: $O(\log n)$

Merge Sort with Parallel Merge:

- ▶ parallel merge: $O(n)$ work and $O(\log^2 n)$ depth
- ▶ work of merge sort still $O(n \log n)$
- ▶ depth: recurrence relation for $d(n) = d(n/2) + O(\log^2 n)$, yields $O(\log^3 n)$
- ▶ parallelism: $O(n / \log^2 n)$

Hillis and Steele Scan

```
for  $k \leftarrow 0$  to  $\log n - 1$  do  
     $j \leftarrow 2^k$   
     $\{a[i] \leftarrow a[i - j] + a[i] : i \in [0..n) \mid i \geq j\}$   
end
```

- ▶ $n - 2^k$ independent tasks (each doing an addition) at each level
- ▶ tasks in one level complete before next level starts
- ▶ depth: $O(\log n)$, work: $O(n \log n)$ (not work efficient), parallelism: $O(n)$

Blelloch Scan

Two traversals of reduction tree, so same complexity as reduction: $O(n)$ work, $O(\log n)$ depth, and $O(n/\log n)$ parallelism

Performance Metrics

- ▶ Speedup
- ▶ Cost
- ▶ Efficiency
- ▶ Throughput

Speedup

- ▶ *relative speedup*: ratio of the execution time of the parallel program on one core to the time on p cores
- ▶ *absolute speedup*: ratio of the execution time of the best known sequential algorithm to the execution time of the parallel algorithm on p cores.

Cost

Cost: product of the asymptotic execution time and the number of cores.

- ▶ Not the same as work!
- ▶ Reduction: $O(n)$ work. If $p = n/2$ cores used, cost is $O(p \log p)$. Reduction is work efficient but is not cost efficient
- ▶ Reduction: if fewer cores used, then each core first sums n/p before parallel reduction. This takes $O(n/p + \log p)$ time, so cost is $O(n + p \log p)$. If n is greater than $p \log p$ then this algorithm is cost efficient and its speedup is $O(n/(n/p)) = O(p)$, which is optimal.

Efficiency

- ▶ *Efficiency*: ratio of the speedup to the number of cores.
- ▶ Also given by ratio of the execution time of the sequential algorithm to the cost

Throughput Metrics

- ▶ When sequential time not possible or meaningful, such as for GPU execution
- ▶ E.g.: floating point operations per second (FLOPS), cell updates per second, ...
- ▶ *Operation count must be based on best known sequential algorithm*

Comparing Scan Algorithms

- ▶ Hillis and Steele with $p = n$ cores: $\log n$ time and $n / \log n$ speedup
- ▶ Blelloch with $p = n/2$ cores: $n / (2 \log n)$ speedup.
- ▶ Both algorithms cost-inefficient, at $O(n \log n)$
- ▶ Better to use fewer cores, as for reduction

SPMD Shared Memory Scan

```
// assumes  $n$  divisible by  $p$   
shared  $a, b$   
 $start \leftarrow id * n/p$   
 $sum \leftarrow 0$   
for  $j \leftarrow 0$  to  $n/p - 1$  do  
     $sum \leftarrow sum + a[start + j]$   
     $a[start + j] \leftarrow sum$   
end  
 $b[id] \leftarrow a[(id + 1) * n/p - 1]$  // sum of all values in  
    sub-array  
 $scan(b)$   
if  $id > 0$  then  
    for  $j \leftarrow 0$  to  $n/p - 1$  do  
         $a[start + j] \leftarrow a[start + j] + b[id - 1]$   
    end  
end
```

Analysis of SPMD Scan


Parallel execution time:

- ▶ Hillis and Steele: $2\lceil n/p \rceil + \log p$
- ▶ Blelloch: $2\lceil n/p \rceil + 2\log p$


Cost is $O(n + p \log p)$, so cost efficient if n is greater than $p \log p$

Speedup of SPMD Hillis and Steele scan

$$n/(2\lceil n/p \rceil + \log p)$$



Efficiency of SPMD Hillis and Steele scan





Strong and Weak Scaling

- ▶ *Strong scaling*: problem size fixed as number of cores increases
- ▶ *Weak scaling*: amount of work per core is kept fixed as number of cores increases

Weak Scaling of SPMD Hillis and Steele

Problem size scaled as $n = 512p$






Communication Analysis

- ▶ Communication over a link: $t_{comm} = \lambda + m/\beta$, for HW/SW latency λ and bandwidth β .
- ▶ Communication topology of application: *virtual topology*
- ▶ At runtime, virtual topology embedded on the physical topology

Communication Analysis: Reduction



$t_{\text{reduction}} = (\sigma + \lambda + m/\beta) \log p = O(m \log p)$, σ is time to add two values and m is size of a value.

Analysis of Game of Life

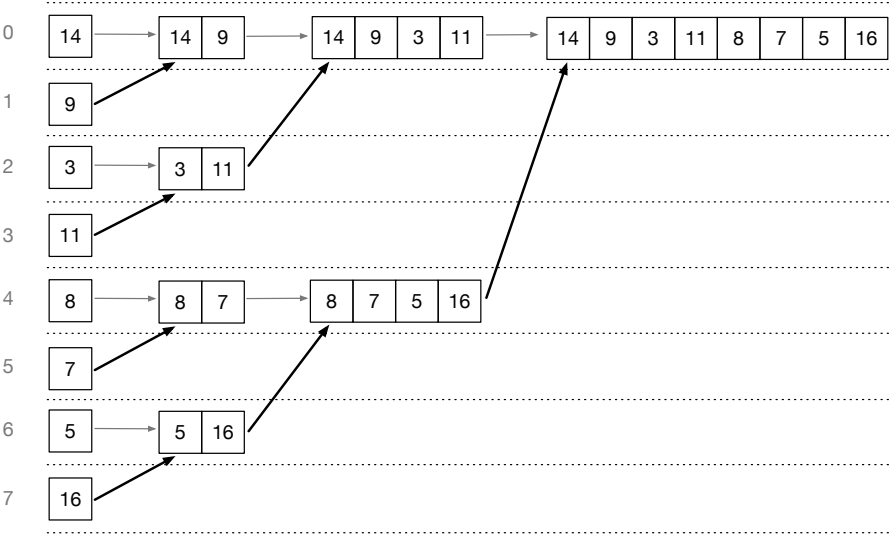
- ▶ 1D decomposition: each core sends 2 messages of size n , $t_{1D} = O(n^2/p + n)$
- ▶ 2D decomposition: each core sends 4 messages of size n/\sqrt{p} , $t_{2D} = O(n^2/p + n/\sqrt{p})$
- ▶ Both decompositions are cost efficient
- ▶ 2D decomposition has lower time and better weak scaling (increasing n as a function of \sqrt{p}), with an efficiency of $O(1)$

Analysis of 1D Matrix-Vector Multiplication

- ▶ communication: gather result vectors and broadcast gathered vector
- ▶ Broadcast has same topology as reduction,
 $t_{\text{broadcast}} = (\lambda + n/\beta) \log p = O(n \log p)$
- ▶ Gather is similar, but message size increases at each step

Gather

core



Gather Analysis

```
// Assumes  $p = 2^i$  and  $n \bmod p = 0$   
for  $i \leftarrow 0$  to  $\log p - 1$  do  
  if  $id \bmod 2^i = 0$  then  
    if  $id/2^i \bmod 2 = 1$  then  
      send array of length  $2^i n/p$  to  $id - 2^i$   
    else  
      receive array starting at position  
       $(id + 2^i)n/p$  from  $id + 2^i$   
    end  
  end  
end
```

- ▶ $\log p$ steps and the message size at step i is $2^i n/p$

1D vs. 2D Matrix-Vector Multiplication

$$t_{1D} = O(n^2/p + n \log p)$$

$$\begin{aligned} t_{2D} &= O(n^2/p + (n/\sqrt{p}) \log \sqrt{p} + (n/\sqrt{p}) \log \sqrt{p}) \\ &= O(n^2/p + (n/\sqrt{p}) \log p) \end{aligned}$$