

Chapter 5: Barriers to Performance

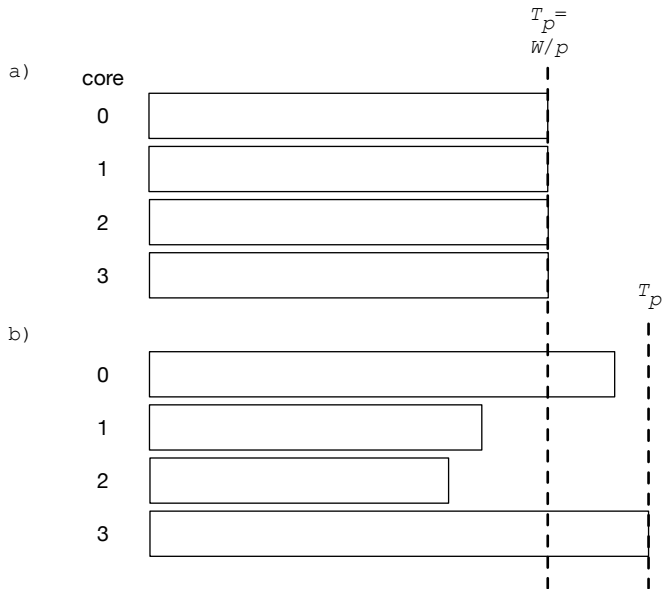
Elements of Parallel Computing

Eric Aubanel

Barriers to Performance

- ▶ Load Imbalance
- ▶ Communication Overhead
- ▶ False Sharing
- ▶ Locality: Hierarchical Algorithms
- ▶ Inherently Sequential Execution
- ▶ Memory
- ▶ Other Overhead

Load Imbalance




Load Balancing

- ▶ Static load balancing
 - ▶ round robin scheduling
- ▶ Dynamic load balancing
 - ▶ Master-worker
 - ▶ work-stealing (fork-join frameworks)

Communication Overhead: Game of Life

Investigate 2 ways to reduce communication overhead:

1. tradeoff redundant communication with reduced communication
2. overlap communication and computation



// each task has $(n/p + 4) \times n$ arrays *grid* and *newGrid*

$nbDown \leftarrow (id + 1) \bmod p$, $nbUp \leftarrow (id - 1 + p) \bmod p$

$m \leftarrow n/p$ // Assume $n \bmod p = 0$

for $k \leftarrow 0$ **to** $N - 1$ **do**

 offset $\leftarrow k \bmod 2$

if offset = 0 **then**

 nonblocking send *grid*[***m*..*m* + 1**, 0..*n* - 1] to *nbDown*

 nonblocking send *grid*[**2..3**, 0..*n* - 1] to *nbUp*

 receive from *nbDown* into

grid[***m* + 2..*m* + 3**, 0..*n* - 1]

 receive from *nbUp* into *grid*[**0..1**, 0..*n* - 1]

end

foreach *cell* at coordinate

$(i, j) \in (1 + \text{offset}..\mathbf{m} + 2 - \text{offset}, 0..n)$ **do**

 updateGridCell(*grid*, *newGrid*, *i*, *j*)

end

 swap references to *newGrid* and *grid*

```

nbDown  $\leftarrow (id + 1) \bmod p$ , nbUp  $\leftarrow (id - 1 + p) \bmod p$ 
m  $\leftarrow n/p$  // Assume  $n \bmod p = 0$ 
for a number of generations do
    nonblocking send grid[m, 0..n - 1] to nbDown
    nonblocking send grid[1, 0..n - 1] to nbUp

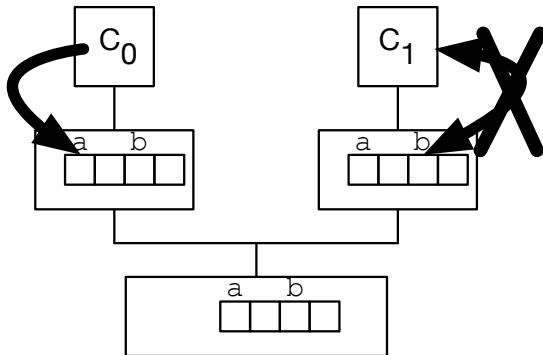
    foreach cell at coordinate  $(i, j) \in (2..m - 1, 0..n)$  do
        updateGridCell(grid, newGrid, i, j)
    end

    // receive boundary values from neighbors
    into ghost elements
    receive from nbDown into grid[m + 1, 0..n - 1]
    receive from nbUp into grid[0, 0..n - 1]
    foreach cell at coordinate
     $(i, j) \in (1, 0..n) \wedge (m, 0..n)$  do
        updateGridCell(grid, newGrid, i, j)
    end

    swap references to newGrid and grid
end

```


False Sharing



- ▶ False sharing of a and b , since intent isn't to share, but sharing is forced because both lie on a cache block.
- ▶ Concurrent access to a and b is not possible.

Locality: Hierarchical Algorithms

- ▶ Challenging to favour local data accesses in hierarchical memory (e.g. GPUs)
- ▶ Hierarchical algorithms can enhance locality
 - ▶ combine shared memory on node and message passing between nodes

Inherently Sequential Execution

Amdahl's law:

$$\text{speedup} = \frac{t_1}{ft_1 + (1 - f)t_1/p}$$

where f is fraction of sequential execution time t_1 for the operations that cannot be done in parallel

- ▶ Not significant problem for properly designed parallel programs
- ▶ Sequential portion often has lower complexity than parallel portion
- ▶ Weak scaling: time of sequential portion decreases as problem size increases

Memory

- ▶ Protecting from data races takes overhead
 - ▶ use private data structures as much as possible
 - ▶ use lock-free algorithms
- ▶ Shared memory communication overhead
 - ▶ ensure locality

Other Overhead

- ▶ Unavoidable
 - ▶ packing data before sending message
- ▶ Avoidable
 - ▶ needlessly allocating memory