


# Chapter 6: Single Source Shortest Path

Elements of Parallel Computing

Eric Aubanel


# Example



# Definitions

- ▶ **Shortest path** between two vertices of a graph: path that minimizes sum of edge weights
- ▶ **SSSP**: find shortest paths between source vertex all other vertices
- ▶ **Diameter** of graph: length of the longest shortest path between all pairs of vertices, ignoring weights
- ▶ **Scale-Free Graphs**: degree distribution follows a power law, that is, the number of vertices of degree  $d$  is  $O(d^{-\lambda})$ , where  $\lambda$  is a small constant.

# Scale-Free Graph



# SSSP Solution with Labelling

**Input:** Graph with vertices  $V$  ( $|V| = n$ ) and edges  $E$  with weights  $C$ , source vertex  $s$ .

**Output:**  $D$ : distance between  $s$  and all other vertices.  $P$ : pointer to predecessor to each vertex in shortest path.

**for**  $i \leftarrow 0$  **to**  $n - 1$  **do**  $D[i] \leftarrow \infty$

$D[s] \leftarrow 0$ ,  $\text{list} \leftarrow \{s\}$

**while**  $\text{list} \neq \emptyset$  **do**

    remove vertex  $i$  from list

    // Relax each out-edge of vertex  $i$

**foreach** edge  $e_{ij} \in E$  **do**

**if**  $D[j] > D[i] + c_{ij}$  **then**

$D[j] \leftarrow D[i] + c_{ij}$ ,  $P[j] \leftarrow i$

**if**  $j \notin \text{list}$  **then**  $\text{list} \leftarrow \text{list} \cup \{j\}$

**end**

**end**


**end**

# SSPP Strategies:

- ▶ *Label-setting*: in each iteration the distance value of a single vertex is permanently set
- ▶ *Label-correcting*: distance value of vertices may be updated multiple times

# Data Structures

- ▶ Adjacency lists
- ▶ Adjacency matrix
- ▶ *Compressed sparse row*: adjacency lists stored contiguously



# Bellman-Ford Algorithm

Standard algorithm:

Initialize  $D$

```
for  $k \leftarrow 1$  to  $n - 1$  do  
    foreach edge  $e_{ij} \in E$  do  
        if  $D[j] > D[i] + c_{ij}$  then  
             $D[j] \leftarrow D[i] + c_{ij}$   
        end  
    end  
end
```

Better to use label correcting algorithm above, with FIFO queue



# Bellman-Ford Example

Queue	vertices updated
0(0)	
1(3), 2(1)	1, 2
2(1), 3(5)	3
3(5), 1(2), 4(6)	1, 4
1(2), 4(6), 5(6)	5
4(6), 5(6), 3(4)	3
5(6), 3(4), 6(8)	6
3(4), 6(8), 7(10)	7
6(8), 7(10), 5(5)	5
7(9), 5(5)	7
5(5)	
6(7)	6
7(8)	7

# Dijkstra's Algorithm

List algorithm, where list is a min-priority queue.

<b>Queue</b>	<b>vertices updated</b>
0(0)	
2(1), 1(3)	1, 2
1(2), 4(6)	1, 4
3(4), 4(6)	3
5(5), 4(6)	5
4(6), 6(7), 7(9)	6, 7
6(7), 7(9)	
7(8)	7

# Complexity

- ▶ Bellman-Ford:  $O(|V||E|)$
- ▶ Dijkstra: Min-priority queue takes  $O(\log |V|)$  to update vertices and remove the minimum vertex, so  $O((|V| + |E|) \log |V|)$  overall.
- ▶ In practice Bellman-Ford has lower runtime than complexity suggests, but no known average case
- ▶ Delta-Stepping algorithm is a compromise between Bellman-Ford and Dijkstra, and provides average-case linear complexity

# Delta-Stepping

- ▶ Vertices placed into array  $B$  of buckets.
- ▶ Bucket  $i$  holds vertices with distance in the range  $[i\Delta, (i + 1)\Delta)$
- ▶ Each relaxation will place a vertex in one of the buckets, and may also move it from another bucket

# Delta=5

<b>B[0]</b>	<b>B[1]</b>	<b>vertices updated</b>
0(0)		
1(3), 2(1)		1, 2
2(1)	3(5)	3
1(2)	3(5), 4(6)	1, 4
3(4)	4(6)	3
	4(6), 5(5)	5
	5(5), 6(8)	6
	6(7), 7(9)	6, 7
	7(8)	7

# Delta Stepping Algorithm

**Input:** Graph with vertices  $V$  ( $|V| = n$ ) and edges  $E$  with weights  $C$ , source vertex  $s$ .

**Output:**  $D$ : distance between  $s$  and all other vertices.

**foreach** vertex  $v \in V$  **do** //classify edges as light or heavy

$H[v] \leftarrow \{e_{vw} \in E \mid c_{vw} > \Delta\}$

$L[v] \leftarrow \{e_{vw} \in E \mid c_{vw} \leq \Delta\}$

$D[v] \leftarrow \infty$

**end**

$\text{relax}(s, 0)$  // places  $s$  in first bucket ( $B[0]$ )

$i \leftarrow 0$

```

while  $B$  is not empty do
     $R \leftarrow \emptyset$ 
    while  $B[i] \neq \emptyset$  do
         $Req \leftarrow \{(w, D[v] + c_{vw}) \mid v \in B[i], e_{vw} \in L[v]\}$ 
         $R \leftarrow R \cup B[i]$  // remember vertices deleted
                           from bucket
         $B[i] \leftarrow \emptyset$ 
        foreach  $(w, x) \in Req$  do
            relax  $(w, x)$ 
        end
    end
     $Req \leftarrow \{(w, D[v] + c_{vw}) \mid v \in R, e_{vw} \in H[v]\}$ 
    foreach  $(w, x) \in Req$  do
        relax  $(w, x)$ 
    end
     $i \leftarrow i + 1$ 
end

```

```
// Relax edge to vertex  $w$  with candidate weight  $x$   
// If accepted assign to appropriate bucket
```

```
Procedure relax( $w, x$ )
```

```
    if  $x < D[w]$  then
```

```
         $B[\lfloor D[w]/\Delta \rfloor] \leftarrow B[\lfloor D[w]/\Delta \rfloor] \setminus \{w\}$ 
```

```
         $B[\lfloor x/\Delta \rfloor] \leftarrow B[\lfloor x/\Delta \rfloor] \cup \{w\}$ 
```

```
         $D[w] \leftarrow x$ 
```

```
    end
```

```
end
```



# Task Decomposition

Bellman Ford, decompose:

- ▶ iterations of foreach loop (edges adjacent to a vertex)
- ▶ process contents of list in parallel



# Dijkstra vs. Bellman Ford

- ▶ *Dijkstra*: only have independent tasks for relaxations of edges from one vertex
- ▶ *Dijkstra*:  $|V| \log |V|$  depth and  $O((|V| + |E|) \log |V|)$  work, so  $O(|E|/|V|)$  parallelism
- ▶ *Bellman-Ford*: in worst case,  $O(|V|)$  depth,  $O(|E||V|)$  work, and  $O(|E|)$  parallelism
- ▶ *Delta-Stepping*: tune  $\Delta$  to balance parallelism and work-efficiency

# Data Decomposition: Edge Partitions

